POLITECNICO
MILANO 1863

Concept-Based Explanations for Image Classifiers Using Textual Prompts

Daniele Di Santi

Academic year: 2024/2025

Advisor: Prof. Marco Brambilla

Co-advisors: Riccardo Campi, Matteo Bianchi, Antonio De Santis

Black-box Systems

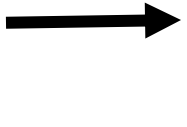

Zebra
p = 0.98

Why is it "Zebra"?
Is there any bias?
Can I trust the model?

Input CNN Prediction User

Explainable Artificial Intelligence (XAI)

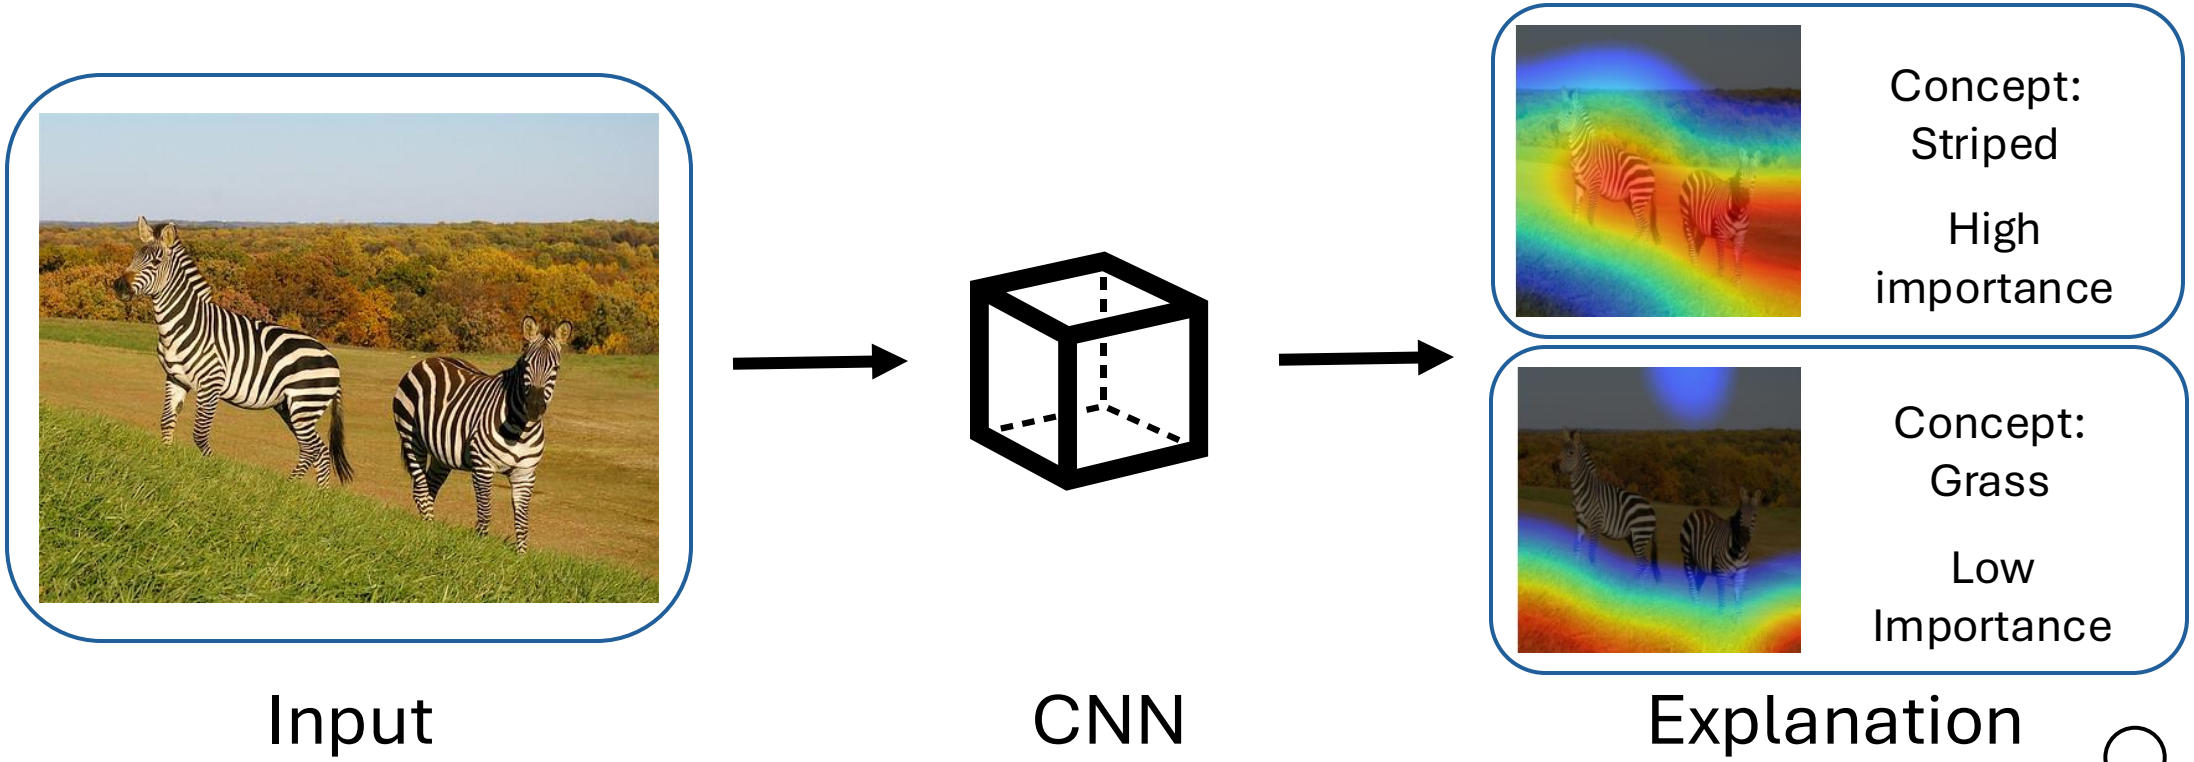

Visual TCAV

Visual TCAV

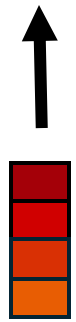

*RESNET*50 architecture
striped concept

*layer*1 layer

Class	Attrib.
zebra	**0.055**
gazelle	**0.026**
impala	**0.027**

*layer*2 layer

Class	Attrib.
zebra	**0.035**
gazelle	**0.019**
impala	**0.017**

*layer*3 layer

Class	Attrib.
zebra	**0.031**
gazelle	**0.014**
impala	**0.012**

*layer*4 layer

Class	Attrib.
zebra	**0.16**
gazelle	**0.039**
impala	**0.032**

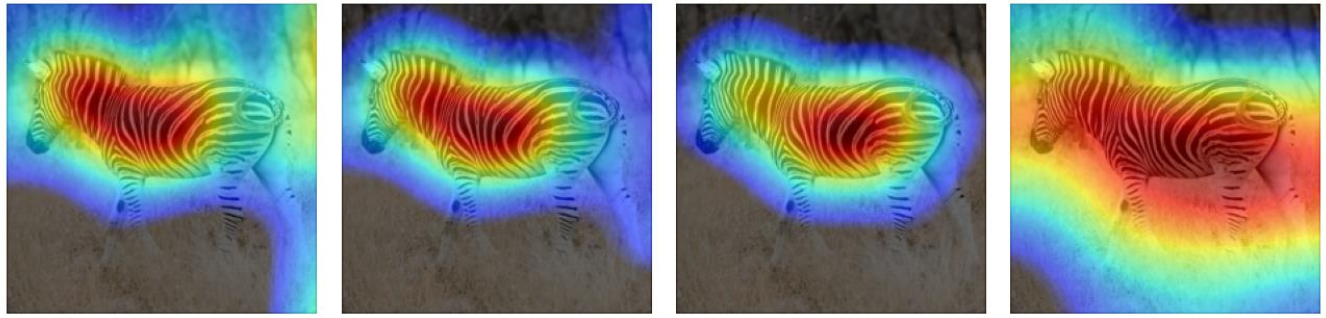

Input image **CAV** **Concept map + concept attributions**

Antonio De Santis et al., Visual-tcav: Concept-based attribution and saliency maps for post-hoc explainability in image classification, 2025.

Problem: Standard CAV Extraction

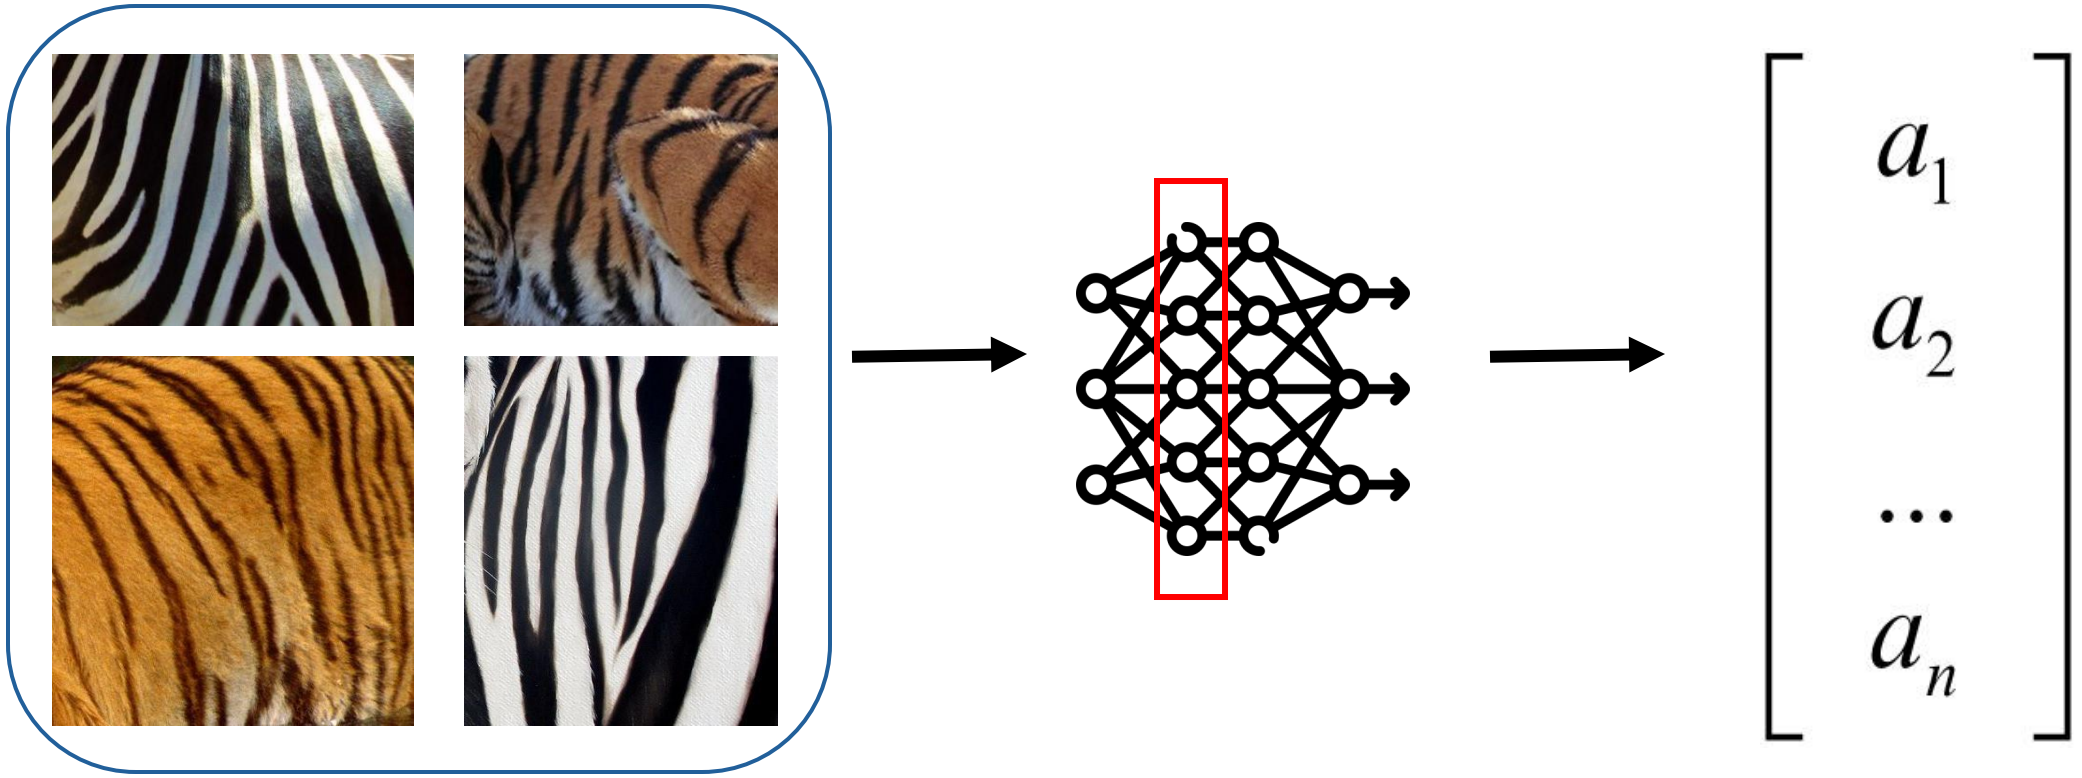

Set of concept images

CNN

CAV

Other Solution: Synthetic Images

"Images of newspaper"

Gen. AI

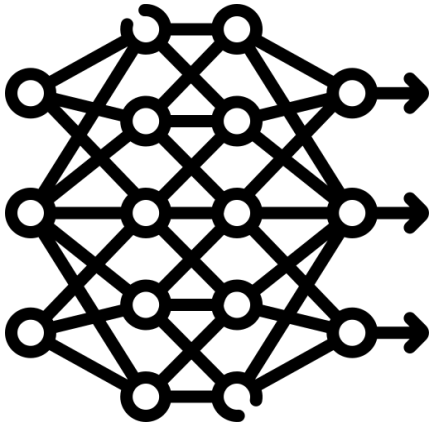

$$\begin{bmatrix} a_1 \\ a_2 \\ \dots \\ a_n \end{bmatrix}$$

Textual Prompt

Set of synthetic concept images

CNN

CAV

Riccardo Campi et al., Towards synthetic concept activation vectors via generative models. In Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR) Workshops, pages 2720–2728, June 2025.

Our idea: from text to CAV

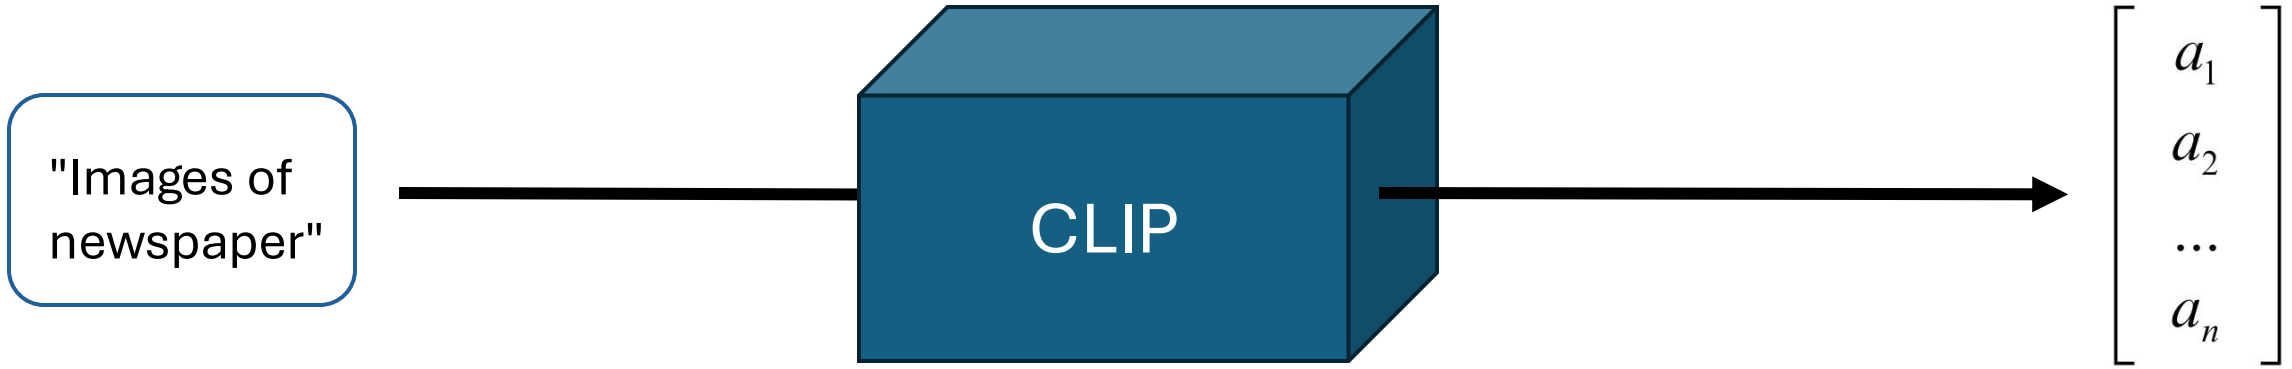

Textual Prompt

Multimodal Model

CAV

Our Solution

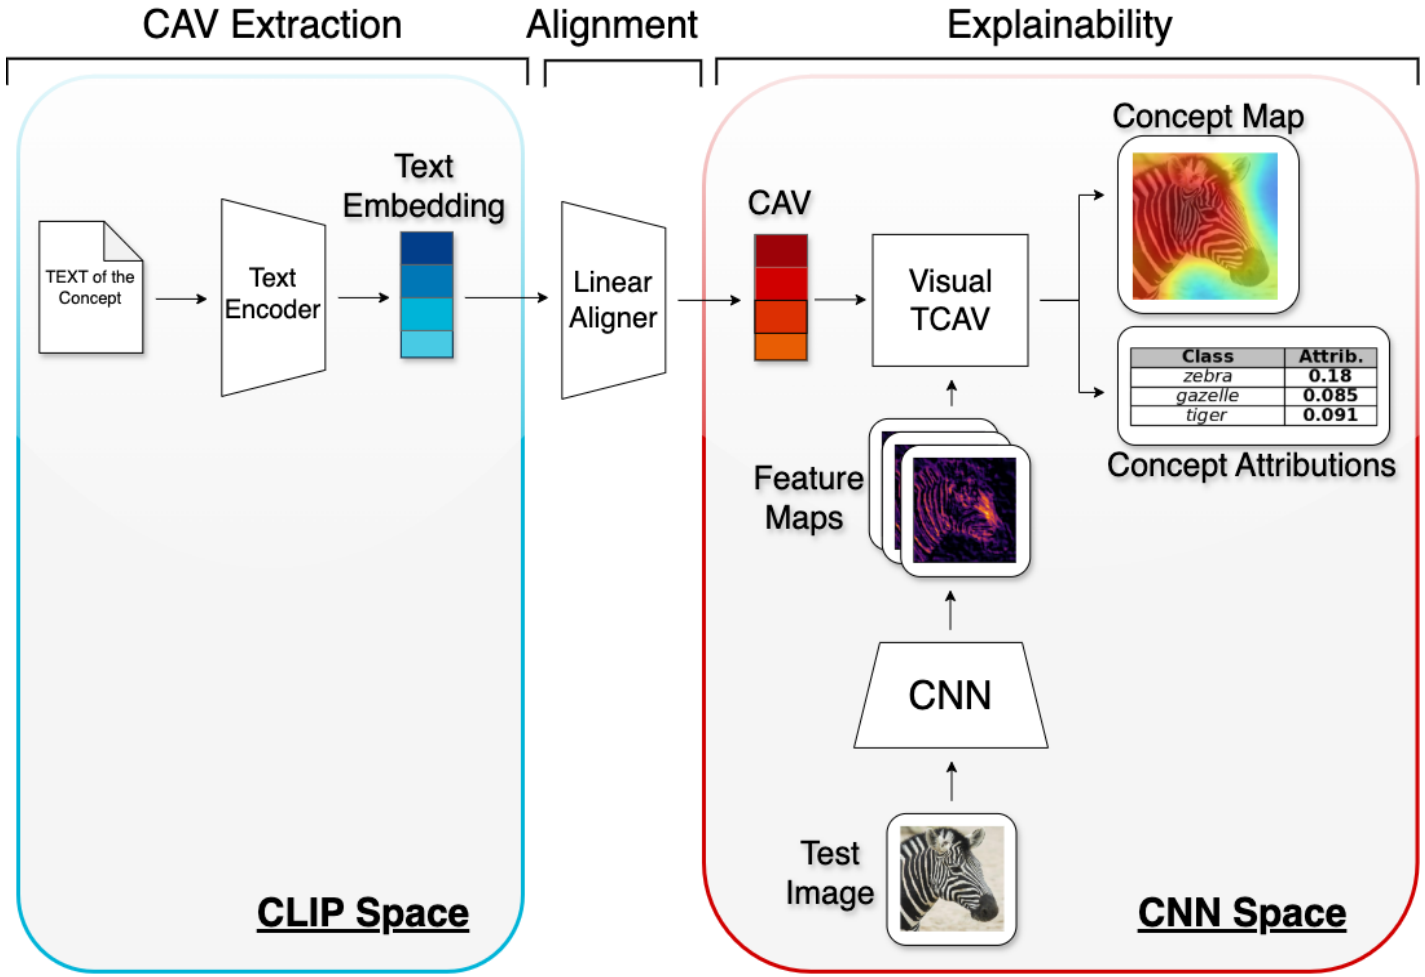

Other use: Localization

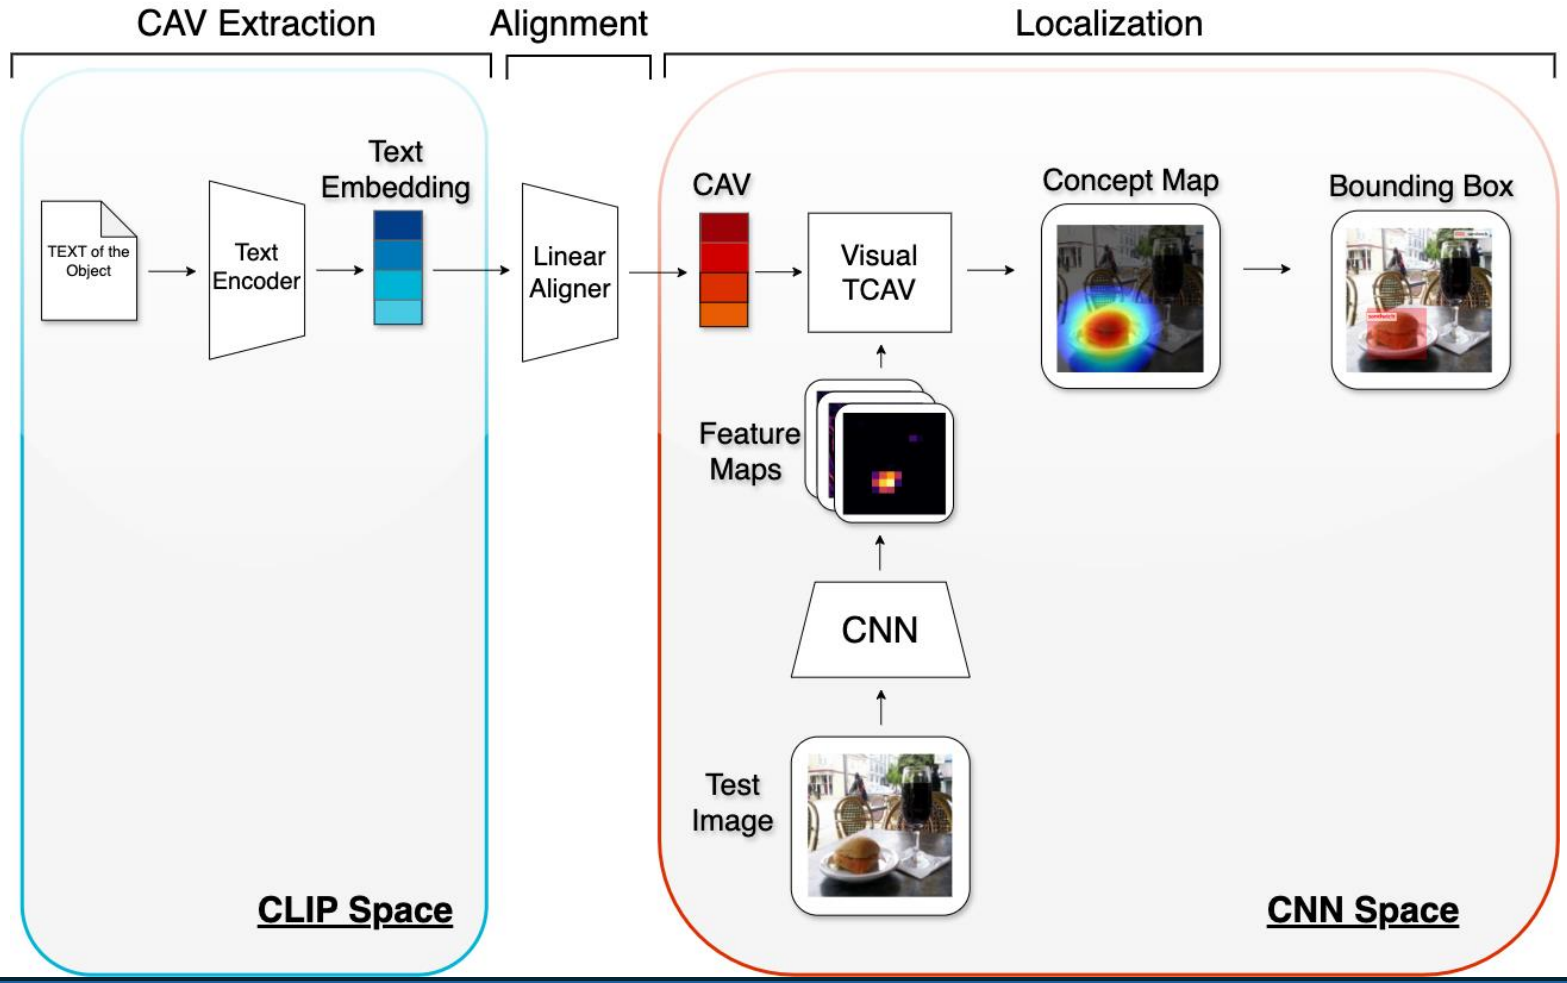

Experiments Setup

PyTorch

Resnet50 pretrained with ImageNet-1k

ImageNet-1k

CONV CONV CONV CONV CONV

Concepts

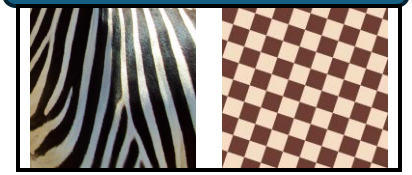

Resnet18 pretrained with CUB-200-2011

CUB-200-2011

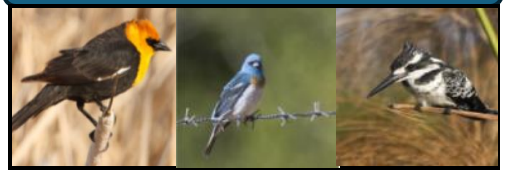

CONV CONV CONV CONV CONV

Concepts

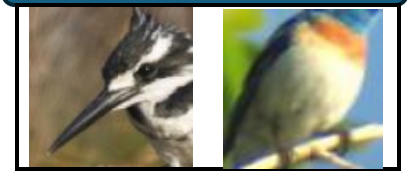

Explanation

Resnet50 pretrained with ImageNet-1k

COCO

CONV CONV CONV CONV CONV

Concepts

Localization

Linear Aligners Training

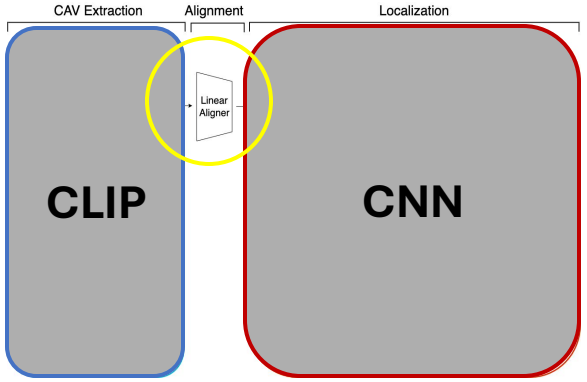

Resnet50 pretrained with ImageNet-1k

Layer	MSE	R^2
layer1	0.492	0.891
layer2	1.14	0.74
layer3	1.798	0.6
layer4	3.205	0.288

Resnet18 pretrained with CUB-200-2011

Layer	MSE	R^2
stage1	0.328	0.927
stage2	0.819	0.818
stage3	1.447	0.678
stage4	2.853	0.366

Experiments: ImageNet-1k

*RESNET*50 architecture
chequered concept

layer4 layer

Class	Attrib.
crossword_pu..	**0.1**
digital_watc..	**0.022**
rule	**0.014**

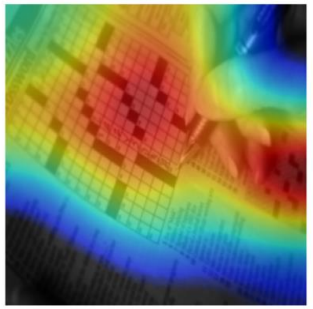

chequered

*RESNET*50 architecture
hand concept

layer4 layer

Class	Attrib.
crossword_pu..	**0.024**
digital_watc..	**0.01**
rule	**0.011**

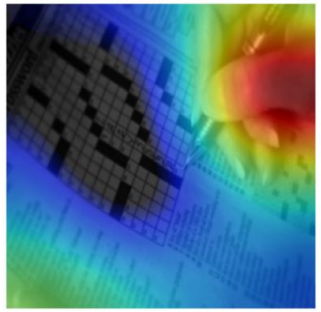

hand

*RESNET*50 architecture
waffled concept

layer4 layer

Class	Attrib.
waffle_iron	**0.13**
mouse	**0.027**
computer_key..	**0.041**

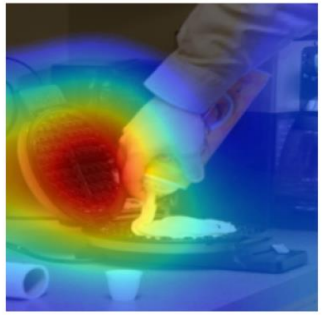

waffled

*RESNET*50 architecture
hand concept

layer4 layer

Class	Attrib.
waffle_iron	**0.012**
mouse	**0.0084**
computer_key..	**0.008**

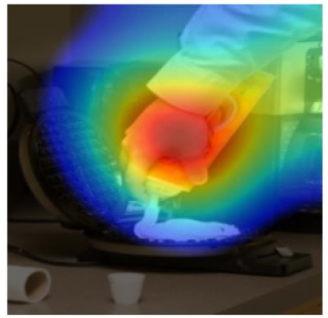

hand

*RESNET*50 architecture
honeycombed concept

layer4 layer

Class	Attrib.
honeycomb	**0.2**
bee	**0.096**
apiary	**0.1**

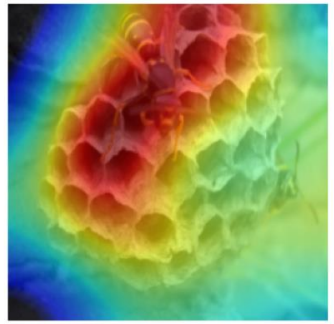

honeycombed

*RESNET*50 architecture
bee concept

layer4 layer

Class	Attrib.
honeycomb	**0.007**
bee	**0.0051**
apiary	**0.0039**

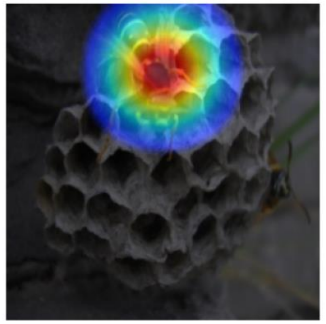

bee

Experiments: ImageNet-1k

*RESNET*50 architecture
legs concept

*layer*1 layer

Class	Attrib.
zebra	**0.29**
gazelle	**0.13**
impala	**0.15**

*layer*2 layer

Class	Attrib.
zebra	**0.035**
gazelle	**0.017**
impala	**0.017**

*layer*3 layer

Class	Attrib.
zebra	**0.047**
gazelle	**0.021**
impala	**0.019**

*layer*4 layer

Class	Attrib.
zebra	**0.082**
gazelle	**0.049**
impala	**0.039**

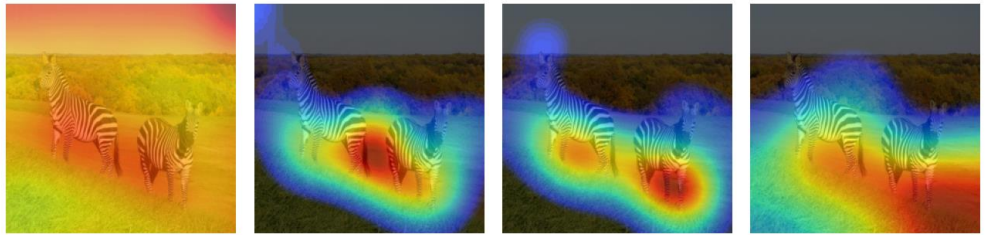

*RESNET*50 architecture
sky concept

*layer*1 layer

Class	Attrib.
zebra	**0.019**
gazelle	**0.011**
impala	**0.015**

*layer*2 layer

Class	Attrib.
zebra	**0.0e+00**
gazelle	**0.0e+00**
impala	**0.0e+00**

*layer*3 layer

Class	Attrib.
zebra	**0.008**
gazelle	**0.0033**
impala	**0.0032**

*layer*4 layer

Class	Attrib.
zebra	**0.0067**
gazelle	**0.0035**
impala	**0.0041**

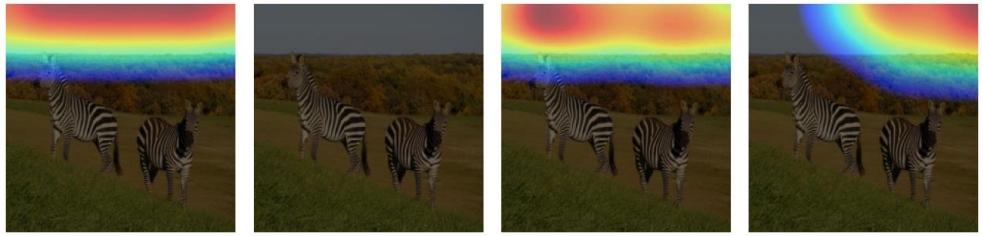

*RESNET*50 architecture
striped concept

*layer*1 layer

Class	Attrib.
zebra	**0.26**
gazelle	**0.12**
impala	**0.14**

*layer*2 layer

Class	Attrib.
zebra	**0.17**
gazelle	**0.09**
impala	**0.095**

*layer*3 layer

Class	Attrib.
zebra	**0.093**
gazelle	**0.032**
impala	**0.027**

*layer*4 layer

Class	Attrib.
zebra	**0.26**
gazelle	**0.055**
impala	**0.04**

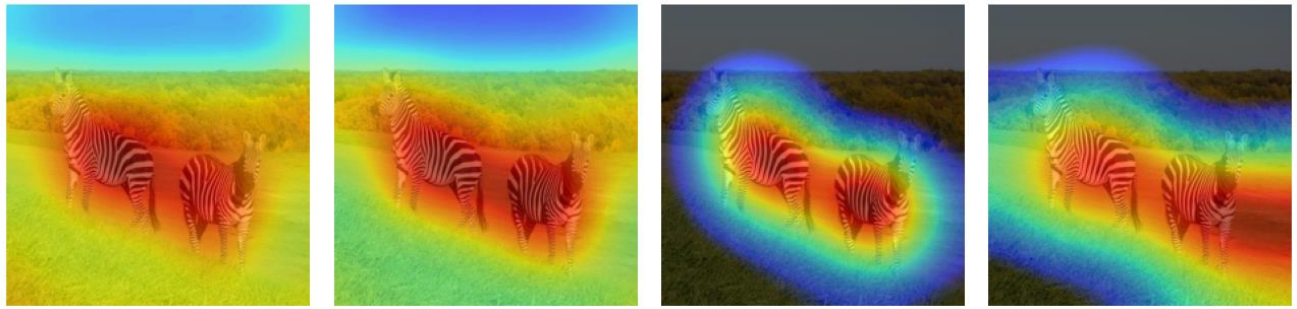

Experiments: ImageNet-1k

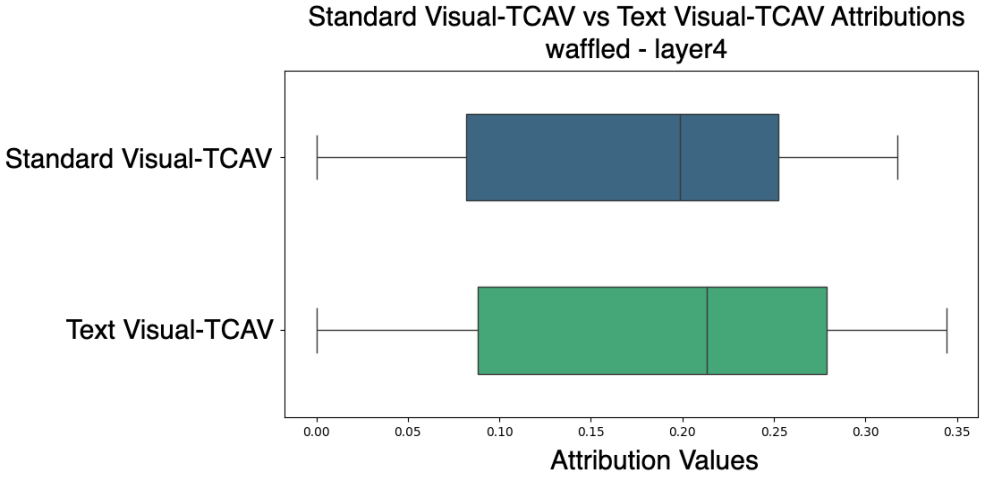

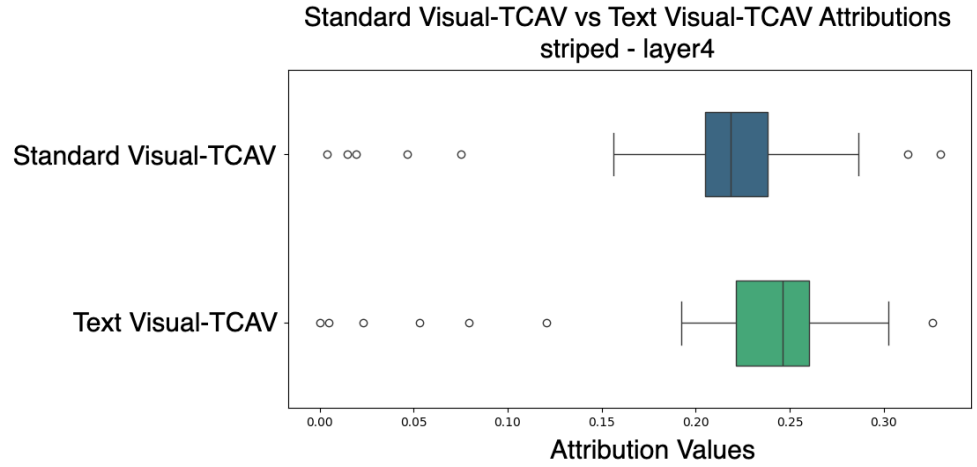

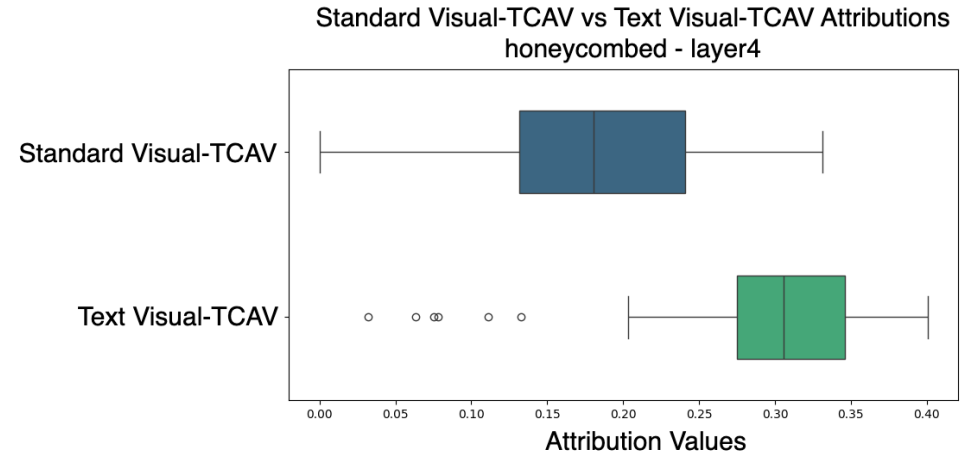

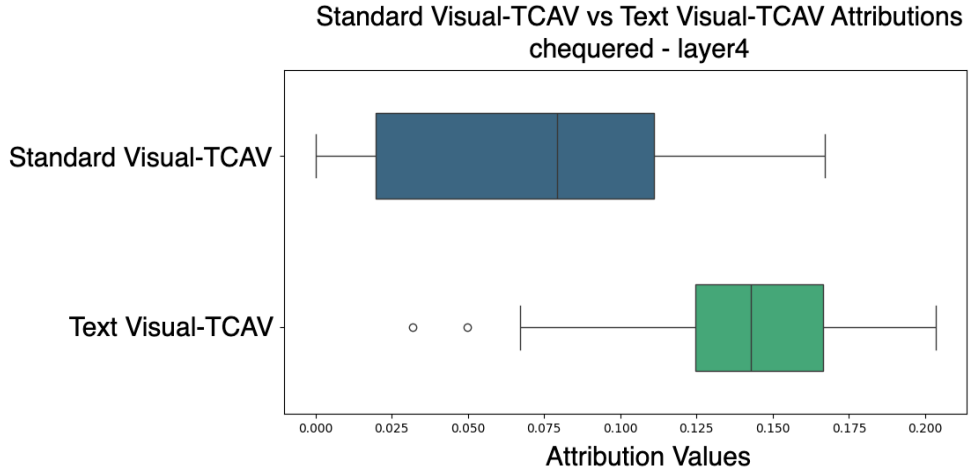

Experiments: CUB-200-2011

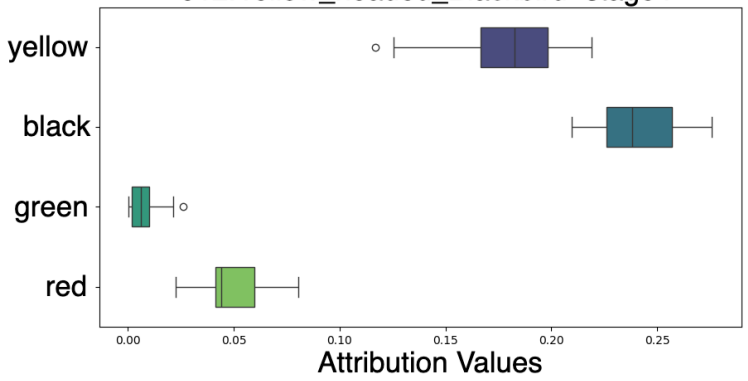

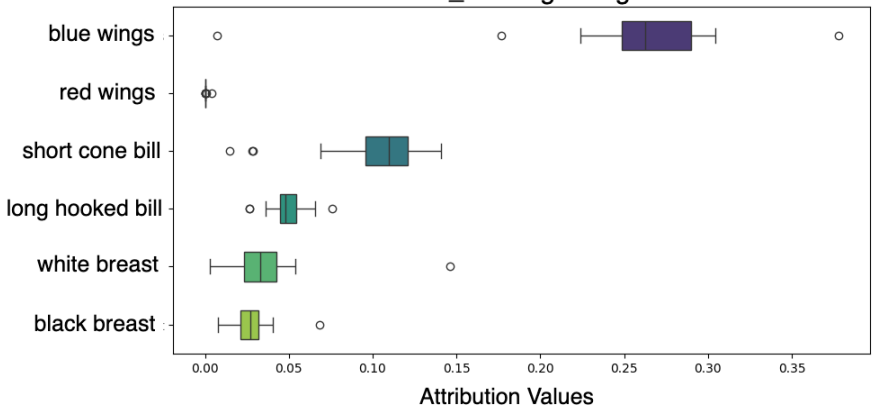

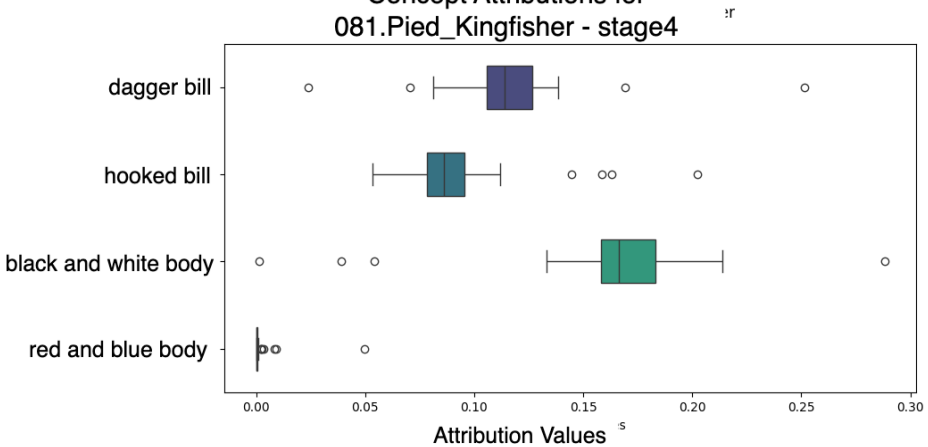

*RESNET*18 architecture
daggerbill concept

*stage*1 layer

Class	Attrib.
Pied_Kingfis..	**0.0e+00**
Black_and_wh..	**0.0e+00**
Green_Kingfi..	**0.0e+00**

*stage*2 layer

Class	Attrib.
Pied_Kingfis..	**0.0e+00**
Black_and_wh..	**0.0e+00**
Green_Kingfi..	**0.0e+00**

*stage*3 layer

Class	Attrib.
Pied_Kingfis..	**0.012**
Black_and_wh..	**0.0056**
Green_Kingfi..	**0.0066**

*stage*4 layer

Class	Attrib.
Pied_Kingfis..	**0.12**
Black_and_wh..	**0.061**
Green_Kingfi..	**0.07**

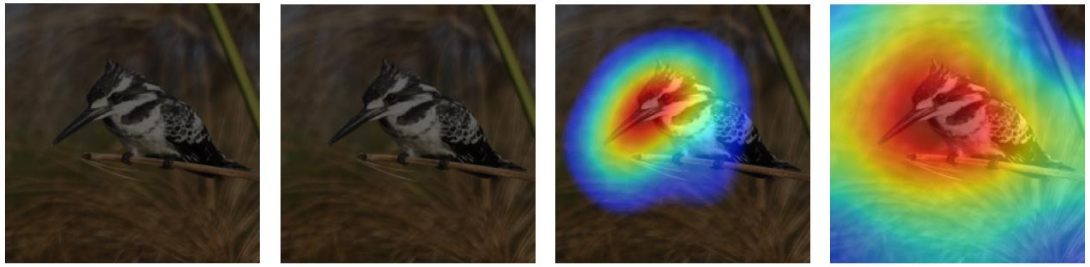

Experiments: Zero-Shot Localization with COCO

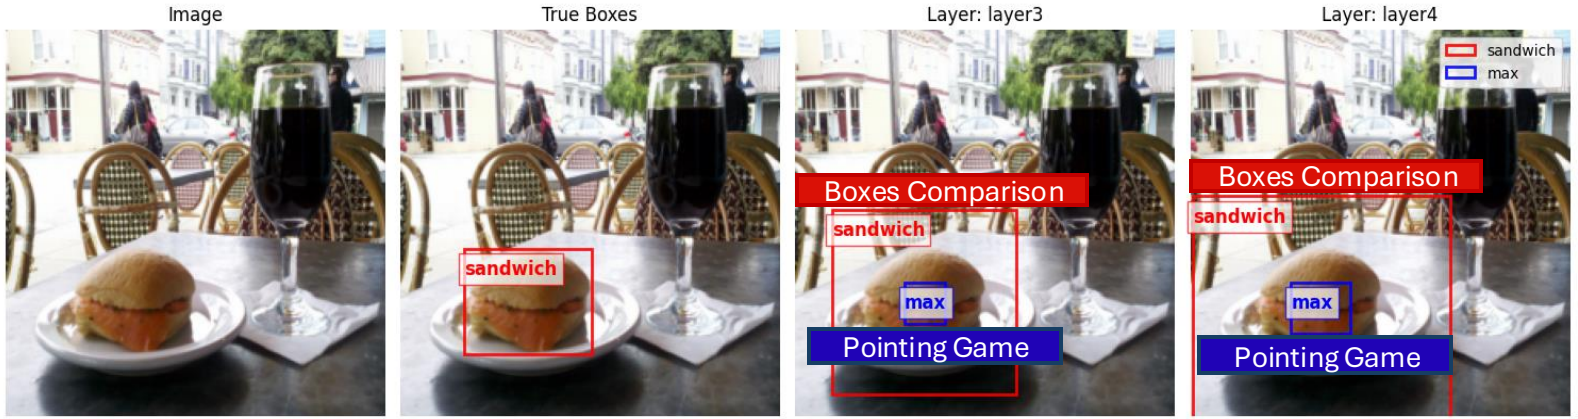

We implement two tasks to assess the performance of zero-shot localization with CAVs from text:

- Pointing Game (correct object position)
- Boxes comparison with IoU (bounding box covers most part of the object)

Accuracy in Pointing Game

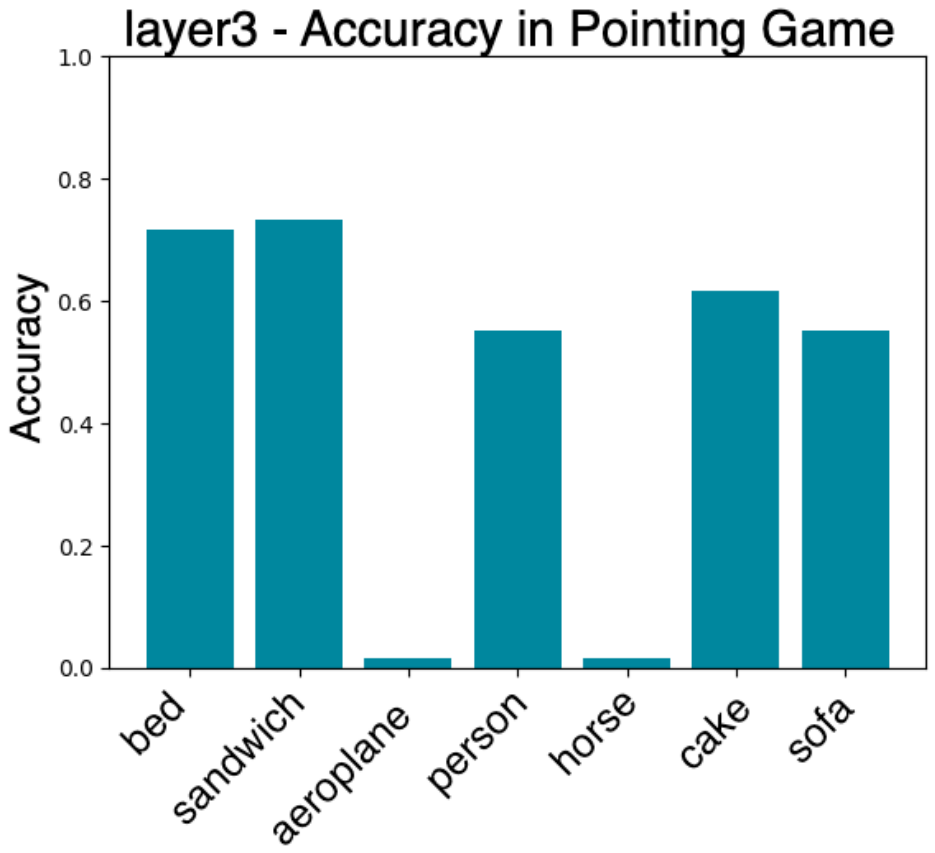

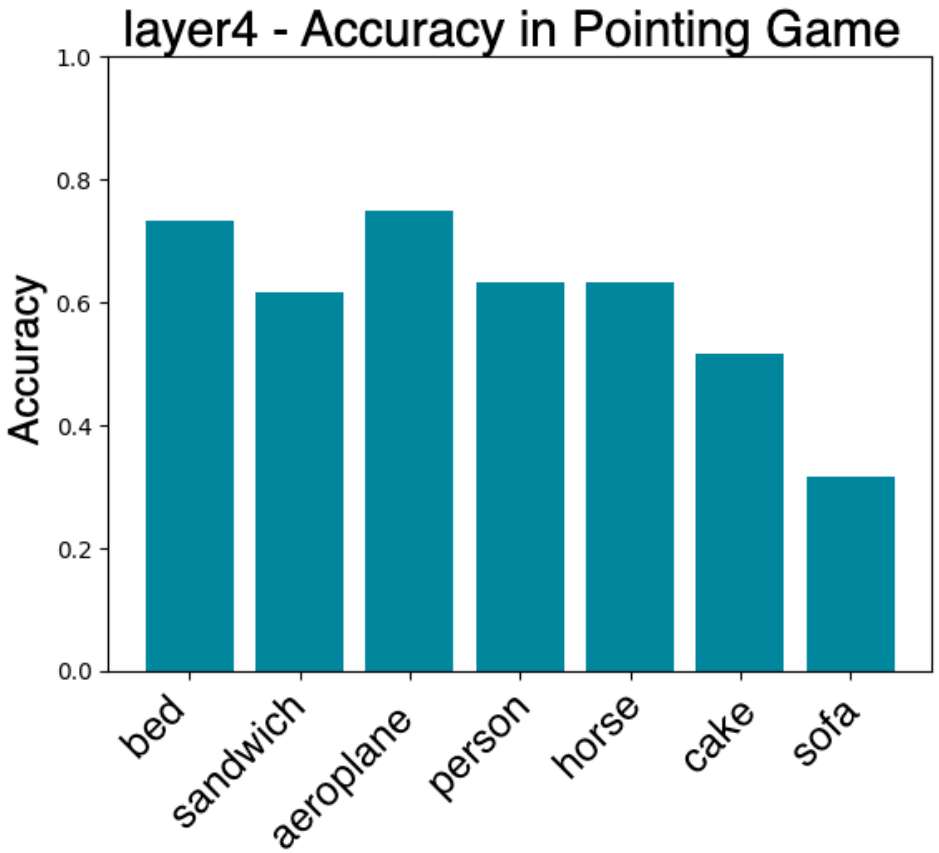

IoU Distribution in Boxes Comparison

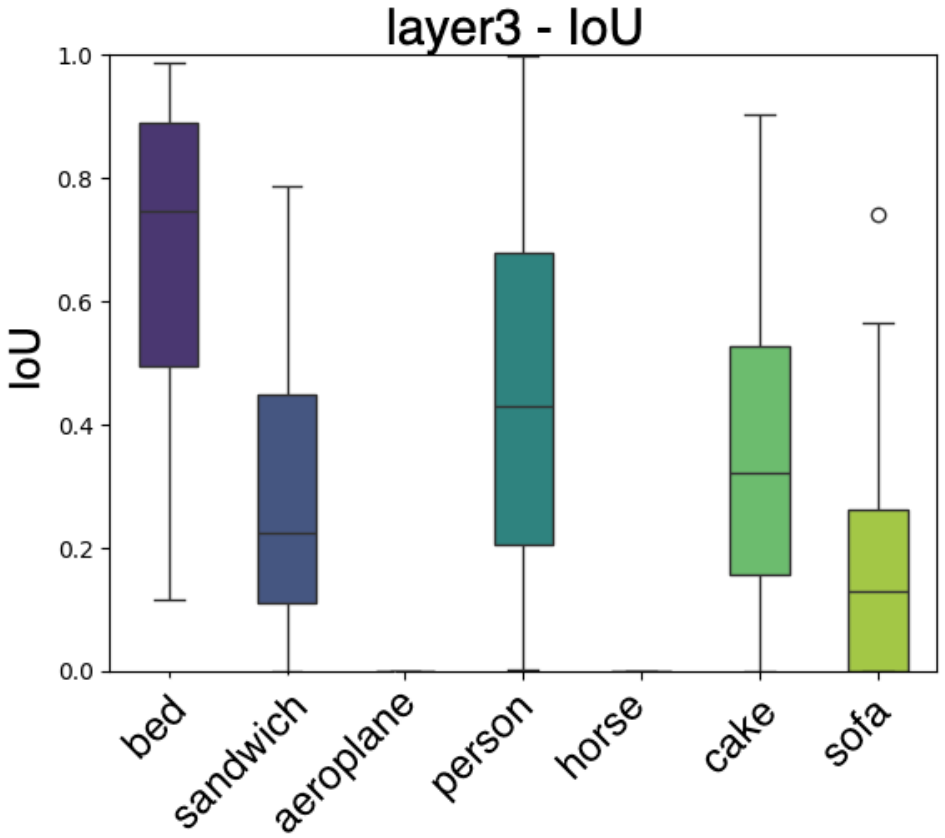

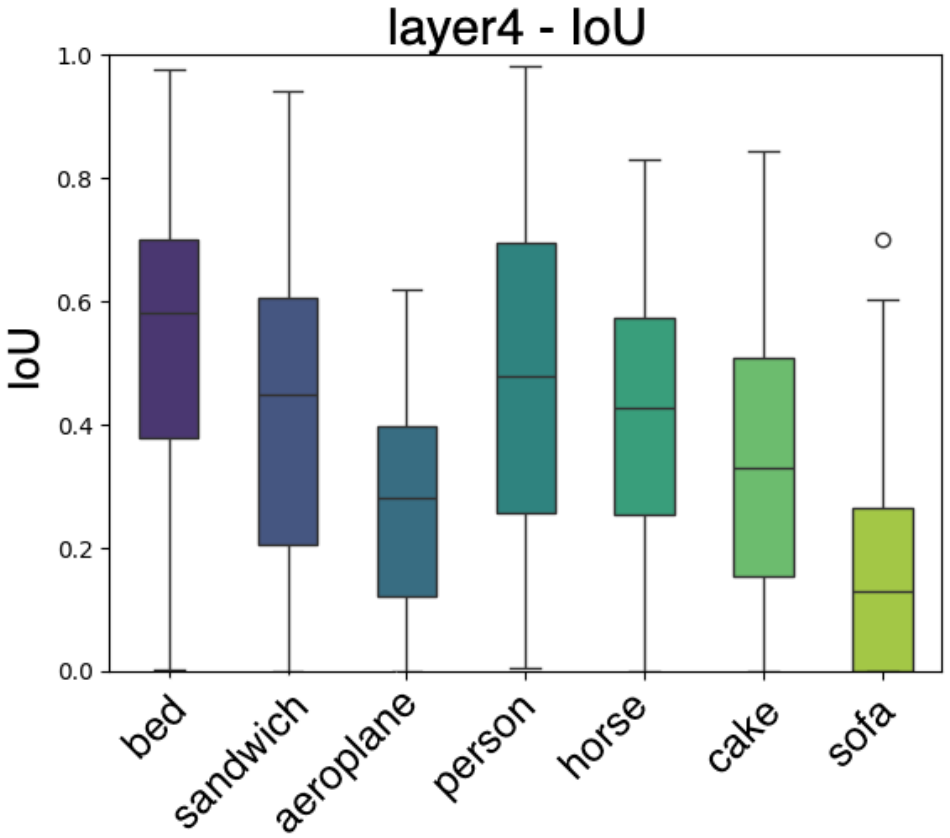

Conclusions

CAVs from text

- **High scalability** for CAV extraction
- **No concept images** for CAV extraction
- **Consistent explanations**
- **Localization** tasks without an annotated dataset
- **Zero-shot localization**

Future work

- More **complex architectures** for the **aligners**
- More roboust ways to **refine** the **concept maps**

POLITECNICO
MILANO 1863

Thank you for your attention